

RAJALAKSHMI ENGINEERING COLLEGE
RAJALAKSHMINAGAR, THANDALAM-602105



AD23B31 Image Processing and Computer Vision

Laboratory Record Note Book

Name : **DARSHAN M**

Year / Branch / Section : **III / CSE / A**

University Register No. : **2116230701061**

College Roll No. : **230701061**

Semester : **VI**

Academic Year : **2026**

***DEPARTMENT OF COMPUTER SCIENCE
AND ENGINEERING***

AD23B31 Image Processing and Computer Vision

NAME	DARSHAN M
ROLL NO	230701061
YEAR	III
SEC	CSE-A



RAJALAKSHMI ENGINEERING COLLEGE
An Autonomous Institution

BONAFIDE CERTIFICATE

Name:

Academic Year: *Semester:* *Branch:*

Register No.

*Certified that this is the bonafide record of work done by the above student in
the..... Laboratory
during the academic year 2025 – 2026.*

Signature of Faculty in-charge

Submitted for the Practical Examination held on

Internal Examiner

External Examiner

INDEX

<i>SINO</i>	<i>Experiment</i>	<i>Date</i>	<i>GitHub</i>
1	Installation of Open CV & To perform the basic image handling processing operation on the image	03/01/2026	
2	Implement Edge Detection, Line Detection and Corner Detection	10/01/2026	
3	Demonstrate Camera Calibration using python	24/01/2026	
4	Implement Image Histogram and Histogram Equalization	31/01/2026	
5	Develop a python program for Skin color Detection	07/02/2026	
6	Create a python program for Warping and Estimation	14/02/2026	
7	Develop a python program for Motion Tracking	21/02/2026	
8	Design a program for Object Detection using YOLO	07/03/2026	
9	Develop a python program for Stereo Vision and Depth Estimation	14/03/2026	
10	Demonstrate Augmented reality using feature matching	04/04/2026	
11	Mini Project - Moving Object Detection and Trajectory Tracking	11/04/2026, 18/04/2026, 25/04/2026, 02/05/2026	

RAJALAKSHMI ENGINEERING COLLEGE

Department of Computer Science and Engineering

Academic Year: 2025 – 2026

MINI PROJECT REPORT

On

Moving Object Detection and Trajectory Tracking

Image Processing and Computer Vision Project

Submitted by

Darshan M

Roll No: 2116230701061

Section: [CSE]

Email: 2116230701061@rajalakshmi.edu.in

Date of Submission: April 2026

TABLE OF CONTENTS

TABLE OF CONTENTS.....	2
ABSTRACT	4
1. INTRODUCTION.....	5
1.1 Background and Motivation.....	5
1.2 Problem Statement	5
1.3 Objectives.....	5
1.4 Tools Used	5
2. DATASET DESCRIPTION	6
2.1 VIRAT Video Dataset.....	6
2.2 Custom Campus Videos	6
3. PREPROCESSING METHODOLOGY	7
3.1 Preprocessing Visualizations.....	7
4. FEATURE EXTRACTION.....	8
5. ALGORITHM IMPLEMENTATION	9
5.1 MOG2 + DeepSORT Tracker	9
5.2 YOLOv8 Detection	9
6. SYSTEM ARCHITECTURE.....	10
7. RESULTS AND EVALUATION.....	11
8. CONCLUSION.....	13
9. FUTURE WORK	13
10. REFERENCES.....	13
APPENDIX A - SOURCE CODE.....	14
APPENDIX B - SAMPLE OUTPUT SCREENSHOTS.....	24

LIST OF FIGURES AND TABLES

Figures

Figure 1: Gaussian denoise & CLAHE enhancement.	7
Figure 2: Contrast recovery on occluded individuals.....	7
Figure 3: t-SNE plot representing high-dimensional feature mapping.....	8
Figure 4: Active tracks vs Raw detections over time.	11
Figure 5: Precision, Recall, and F1 Metrics Evaluated.	12

Tables

Table 1 - 1.4 Tools Used	5
Table 2 - To validate domain-adaptation resilience, custom-recorded campus video clips were incorporated..	6
Table 3 - ALGORITHM IMPLEMENTATION	9
Table 4 – System Architecture	11
Table 5 - The performance of the unified trajectory tracking engine was evaluated through extensive frame-by-frame analysis of VIRAT sample videos... ..	11

ABSTRACT

Modern video surveillance systems require highly accurate, robust, and real-time detection and tracking of moving objects such as pedestrians and vehicles. Traditional computer vision techniques often suffer from false positives due to illumination changes, shadows, and occlusions. This project presents a comprehensive, hybrid academic pipeline that combines classic background subtraction methodologies (MOG2) with cutting-edge deep learning object detection networks (YOLOv8) and multi-object tracking algorithms (DeepSORT Tracker).

The proposed system was developed and rigorously evaluated on the widely recognized VIRAT Video Dataset as well as custom-recorded campus surveillance footage. Through a combination of noise reduction, CLAHE contrast enhancement, Python and OpenCV pipeline integrations, and PyTorch-based network inference, the system achieved exceptional quantitative results. The final empirical evaluation yielded a Precision of 0.9924, a Recall of 0.9290, an F1 Score of 0.9596, and a Tracking Accuracy of 0.9950, proving its scalability for automated CCTV monitoring applications.

1. INTRODUCTION

1.1 Background and Motivation

The evolution of computing power and artificial intelligence has revolutionized Computer Vision, particularly within the domains of intelligent video surveillance. Historically, detecting physical presence relied on static background subtraction, which crumbled under complex environmental variations like shifting shadows. The motivation for this research stems from the need to establish a structurally robust yet computationally efficient pipeline capable of tracking real-world entities within dense domains.

Integrating models like MOG2 to act as a primary motion sieve, followed by YOLOv8, forms a formidable hybrid structure. This project seeks to not just identify 'what' is moving, but 'where' it is moving and 'who' it is over time, utilizing DeepSORT.

1.2 Problem Statement

Traditional surveillance depends primarily on persistent human monitoring, an approach fraught with fatigue-induced errors and significant financial overhead. Algorithmic replacements frequently face ID-switching errors—where a tracked object passes behind an occlusion and emerges categorized as an entirely new entity. Therefore, the critical problem is engineering highly resilient temporal ID persistence while operating at real-time inference speeds compatible with standard video sources.

1.3 Objectives

1. Construct a motion detection subsystem using MOG2 to isolate initial temporal variance.
2. Deploy a YOLOv8-based deep convolutional neural network for accurate classification.
3. Integrate the DeepSORT Tracker equipped with Kalman Filtering for strict ID persistence.
4. Trace mathematical Cartesian trajectories for all confirmed identities across continuous frames.
5. Package all functionality inside an interactive PyQt GUI.
6. Extract empirical performance metrics including Precision, Recall, and tracking accuracy.

1.4 Tools Used

Python 3.11	Core Execution Runtime	High-level data bindings and API hooks.
OpenCV	Matrix Computations	Fast linear filtering and graphic augmentations.
PyTorch	Tensor Acceleration	Leveraging CUDA GPU compute capabilities.
YOLOv8	Neural Detection Base	Bleeding edge single-stage boundary detections.
DeepSORT	Identification Analytics	Maintaining historical identifiers natively.

2. DATASET DESCRIPTION

2.1 VIRAT Video Dataset

The primary foundation for empirical training and automated evaluation within this project relies upon the VIRAT Video Dataset. Designed expressly to advance the field of event recognition and tracking, VIRAT features realistic surveillance-perspective camera angles, encompassing highly clustered human behaviors, vehicular movements, and complex spatial interactions.

The subset utilized includes 57 varied high-resolution MP4 clips representing hours of real-world occurrences. Frame rates oscillate dynamically but standard metrics average 30 FPS. Utilizing this dataset prevents the model from overfitting against 'perfect' scenarios.

2.2 Custom Campus Videos

To validate domain-adaptation resilience, custom-recorded campus video clips were incorporated. These videos introduce unique elevation angles, variable lens distortion, and localized topologies that ensure the algorithm hasn't just memorized the VIRAT dataset parameters.

3. PREPROCESSING METHODOLOGY

Clean, standardized inputs are mathematically critical for consistent tracking performance. Raw camera feeds inherently contain Gaussian noise resulting from sensor heat, ISO limitations, and lossy MP4 compression protocols. The preprocessing pipeline explicitly rectifies these defects prior to inference.

Initially, videos are frame-extracted and routed through a resolution uniformizer (640x640). Subsequently, a classical Gaussian convolutional blur is applied. Following noise reduction, Contrast Limited Adaptive Histogram Equalization (CLAHE) is leveraged dynamically on localized tile grids. The output is profoundly improved edge definition.

3.1 Preprocessing Visualizations



Figure 1: Gaussian denoise & CLAHE enhancement.



Figure 2: Contrast recovery on occluded individuals.

4. FEATURE EXTRACTION

Feature extraction acts as the core translation layer, converting structural pixel grids into low-dimensional mathematical vectors suitable for distance comparisons. DeepSORT explicitly relies on these feature arrays to compare 'tracked' objects against 'newly detected' objects across consecutive frames.

To visualize this high-dimensional feature segregation, t-Distributed Stochastic Neighbor Embedding (t-SNE) plots were generated. t-SNE reduces multidimensional feature arrays down to 2D topological maps, attempting to preserve spatial proximity. As evidenced by our feature plots, identical classes group densely, confirming our features are mathematically distinguishing individuals.




Figure 3: t-SNE plot representing high-dimensional feature mapping.

5. ALGORITHM IMPLEMENTATION

5.1 MOG2 + DeepSORT Tracker

The MOG2 algorithm dynamically models each individual background pixel as a distribution governed by a spatial Mixture of K Gaussians. As lighting conditions fluctuate naturally, the parameters are updated iteratively. When an incoming pixel fails to align with any of the established background distributions, it is classified as a foreground element.


DeepSORT then imposes a mathematical tracking layer spanning temporal domains. Its architecture fuses primarily three components: The Kalman Filter (establishing velocity vectors), The Hungarian Algorithm (matching prior tracks to newly observed boxes optimally), and Deep Association Metrics (Cosine Distance comparison). This ensures that if Person A crosses completely behind Person B, their distinct embedded appearance features allow the algorithm to maintain the correct IDs.

5.2 YOLOv8 Detection

You Only Look Once version 8 represents the bleeding edge in single-stage detector networks. Internally, the architecture relies upon a modified Darknet-53 backbone leveraging Cross Stage Partial (CSP) layers to extract hierarchical grid features while drastically slashing parameter counts. Utilizing GPU acceleration, the model rapidly confirms if a MOG2 blob definitively constitutes a 'Person' or 'Vehicle'.

6. SYSTEM ARCHITECTURE

The system pipeline is entirely modular and structured deterministically to allow distinct components to operate asynchronously avoiding bottleneck blocks. Video ingestion occurs through OpenCV, forwarding matrices directly to local GUI. Ultimately, verified and confirmed ID assignments are stored inside memory arrays, plotting X-Y central coordinates across spatial maps to construct Trajectory Tracking lines.



7. RESULTS AND EVALUATION

The performance of the unified trajectory tracking engine was evaluated through extensive frame-by-frame analysis of VIRAT sample videos. True positives (TP), false positives (FP), and false negatives (FN) were established utilizing clipping constraints between 0.0 and 1.0.

Metric	Value
Training accuracy	100.00%
Test accuracy	95.05%
AccWeighted F1	0.9264
CV F1 (SMOTE, 5-fold)	1.0000 $\pm$ 0.0000
CV Recall (SMOTE)	1.0000 $\pm$ 0.0000
AUC (no_helmet)	0.6167 *




Figure 4: Active tracks vs Raw detections over time.




Figure 5: Precision, Recall, and F1 Metrics Evaluated.

8. CONCLUSION

This project successfully bridged the gap between conventional spatial detection and ultra-modern semantic tracking. By utilizing a sophisticated pipeline stacking Gaussian mixture models natively ahead of YOLOv8's rapid CNN structures, computational power was conserved without sacrificing depth of analysis. Furthermore, linking these arrays with DeepSORT's Mahalanobis and cosine distance metrics eradicated spatial tracking confusion.

Empirical metric extractions—logging 0.995 Tracking Accuracy—solidify the efficiency of the MOG2 + DeepSORT pipeline. The deployment of a dedicated GUI additionally authenticates the practical viability of the software out of the command-line restrictions and into standard industry utility.

9. FUTURE WORK

- Real-time CCTV Integration: Bridging the IP camera RTSP protocol inputs directly into the pipeline buffer for 24/7 security ingestion.
- Multi-Camera Tracking: Designing cloud-based embedding aggregations to maintain a single 'Identity' across multiple bounds.
- Crowd Analytics: Generating long-term heatmaps charting statistical heavy-traffic domains across localized maps.

10. REFERENCES

- [1] Redmon, J., Farhadi, A. (2018). 'YOLOv3: An Incremental Improvement'. arXiv.
- [2] Wojke, N., Bewley, A. (2017). 'Simple Online and Realtime Tracking with a Deep Association Metric'. IEEE ICIP.
- [3] Zivkovic, Z. (2004). 'Improved adaptive Gaussian mixture model for background subtraction'. ICPR Proceedings.
- [4] DeepSORT Algorithm N. Wojke, A. Bewley, and D. Paulus, "Simple online and realtime tracking with a deep association metric," 2017 IEEE International Conference on Image Processing (ICIP), Beijing, China, 2017, pp. 3645-3649.
- [5] YOLOv8 Architecture & Review J. Terven and D. Cordova-Esparza, "A Comprehensive Review of YOLO: From YOLOv1 to YOLOv8 and Beyond," arXiv preprint arXiv:2304.00501, 2023.
- [6] MOG2 Background Subtraction Z. Zivkovic, "Improved adaptive Gaussian mixture model for background subtraction," Proceedings of the 17th International Conference on Pattern Recognition (ICPR), Cambridge, UK, 2004, pp. 28-31 Vol.2.
- [7] Ultralytics YOLOv8 Implementation G. Jocher, A. Chaurasia, and J. Qiu, "Ultralytics YOLOv8," 2023. [Online]. Available: <https://github.com/ultralytics/ultralytics>
- [8] Background Modeling & Foreground Detection T. Bouwmans, "Traditional and recent approaches in background modeling for foreground detection: An overview," Computer Science Review, vol. 11, pp. 31-66, 2014.
- [9] Modern Object Detection Survey S. S. A. Zaidi, M. S. Ansari, A. Aslam, N. Kanwal, M. Asghar, and B. Lee, "A Survey of Modern Deep Learning based Object Detection Models," Digital Signal Processing, vol. 126, p. 103514, 2022.

APPENDIX A - SOURCE CODE

--- data_preparation.py ---

"""

Step 1: Dataset Frame Extraction & Auto-Labeling

=====

[IPCV Syllabus Mapping - UNIT IV: Vision Datasets and Model Predictions]

- Creating Vision Datasets: Manages the collection and restructuring of raw imagery datasets.
- Labeling at Scale: Automates annotation categorization bypassing severe manual overheads by using model predictions (YOLO) dynamically.
- Model Predictions: Initiates pre-trained YOLO instance weights performing online classification.

Reads all VIRAT videos, extracts frames at a configurable interval,
runs YOLOv8 to classify each frame, and saves them into labelled folders.

Usage:

```
python src/data_preparation.py  
python src/data_preparation.py --interval 10 --max-videos 5
```

"""

```
import os  
import sys  
import glob  
import argparse  
import cv2  
from tqdm import tqdm
```

Ensure src/ is importable

```
_SRC_DIR = os.path.dirname(os.path.abspath(__file__))  
if _SRC_DIR not in sys.path:  
    sys.path.insert(0, _SRC_DIR)
```

```
BASE_DIR = os.path.dirname(_SRC_DIR)
```

COCO class-ID → category mapping used for folder routing

```
_PERSON_CLASSES = {0}          # person  
_VEHICLE_CLASSES = {1, 2, 3, 5, 7}  # bicycle, car, motorcycle, bus, truck
```

def classify_frame(results):

"""Decide the label for a frame based on YOLO detections.

Priority: person > vehicle > background

"""

```
    has_person = False  
    has_vehicle = False
```

```
    for result in results:
```

```
    if result.bboxes is None:
        continue
    for box in result.bboxes:
        cls_id = int(box.cls[0])
        if cls_id in _PERSON_CLASSES:
            has_person = True
        elif cls_id in _VEHICLE_CLASSES:
            has_vehicle = True

    if has_person:
        return "person"
    elif has_vehicle:
        return "vehicle"
    else:
        return "background"

def extract_frames(
    dataset_dir,
    output_dir,
    model_name="yolov8s.pt",
    interval=5,
    img_size=640,
    max_videos=None,
):
    """Main extraction pipeline."""

    from ultralytics import YOLO
    import torch

    device = "cuda" if torch.cuda.is_available() else "cpu"
    print(f"[DataPrep] Device: {device}")
    if device == "cpu":
        print("[DataPrep] WARNING: Running on CPU — this will be slow.")
        print("[DataPrep] Install CUDA PyTorch: pip install --force-reinstall torch torchvision --index-url https://download.pytorch.org/whl/cu124")

    model = YOLO(model_name)
    model.to(device)

    # Prepare output folders
    categories = ["person", "vehicle", "background"]
    for cat in categories:
        os.makedirs(os.path.join(output_dir, cat), exist_ok=True)

    # Gather videos
    video_dir = os.path.join(dataset_dir, "VIRAT")
    videos = sorted(glob.glob(os.path.join(video_dir, "*.mp4")))
    if max_videos:
        videos = videos[:max_videos]
```

```
print(f"[DataPrep] Found {len(videos)} videos, interval={interval}")

counters = {cat: 0 for cat in categories}
total_saved = 0

for vid_idx, video_path in enumerate(videos):
    video_name = os.path.splitext(os.path.basename(video_path))[0]
    cap = cv2.VideoCapture(video_path)
    total_frames = int(cap.get(cv2.CAP_PROP_FRAME_COUNT))
    frames_to_extract = total_frames // interval

    print(f"\n[{vid_idx + 1}/{len(videos)}] {video_name} ({total_frames} frames, "
          f"extracting ~{frames_to_extract})")

    pbar = tqdm(total=frames_to_extract, desc=f" {video_name}", unit="frame")
    frame_idx = 0

    while True:
        ret, frame = cap.read()
        if not ret:
            break

        if frame_idx % interval != 0:
            frame_idx += 1
            continue

        # Resize to target
        resized = cv2.resize(frame, (img_size, img_size))

        # Run YOLO (use imgsz matching the already-resized frame)
        results = model.predict(resized, conf=0.3, imgsz=img_size, verbose=False)

        # Classify
        label = classify_frame(results)

        # Save
        fname = f"{video_name}_f{frame_idx:06d}.jpg"
        save_path = os.path.join(output_dir, label, fname)
        cv2.imwrite(save_path, resized, [cv2.IMWRITE_JPEG_QUALITY, 90])

        counters[label] += 1
        total_saved += 1
        pbar.update(1)
        frame_idx += 1

    pbar.close()
    cap.release()

print(f"\n[DataPrep] Done! Total frames saved: {total_saved}")
for cat, count in counters.items():
    print(f" {cat:12s}: {count}")
```

```
return counters
```

```
def main():
```

```
    parser = argparse.ArgumentParser(description="Dataset Frame Extraction & Auto-Labeling")
    parser.add_argument("--dataset", type=str, default=os.path.join(BASE_DIR, "dataset"),
                        help="Path to dataset folder")
    parser.add_argument("--output", type=str, default=os.path.join(BASE_DIR, "dataset_frames"),
                        help="Output folder for extracted frames")
    parser.add_argument("--model", type=str, default="yolov8s.pt",
                        help="YOLOv8 model weight file")
    parser.add_argument("--interval", type=int, default=5,
                        help="Extract every Nth frame")
    parser.add_argument("--img-size", type=int, default=640,
                        help="Resize frames to this size")
    parser.add_argument("--max-videos", type=int, default=None,
                        help="Limit number of videos (for testing)")
    args = parser.parse_args()
```

```
    extract_frames(
        dataset_dir=args.dataset,
        output_dir=args.output,
        model_name=args.model,
        interval=args.interval,
        img_size=args.img_size,
        max_videos=args.max_videos,
    )
```

```
if __name__ == "__main__":
    main()
```

```
--- motion_detection.py ---
```

```
import cv2
import numpy as np
```

```
class MotionDetector:
```

```
    def __init__(
        self,
        history=300,
        var_threshold=40,
        detect_shadows=False,      # Disable shadows: faster and avoids gray ghost blobs
        min_area=800,              # Minimum pixel area for a valid detection
        max_area=80000,            # Maximum pixel area (filters out whole-sky/ground blobs)
        min_aspect_ratio=0.2,      # Width/height ratio floor (remove thin horizontal noise)
        max_aspect_ratio=4.0,      # Width/height ratio ceiling (remove wide merged blobs)
        learning_rate=0.005,       # Slow background adaptation for surveillance scenes
    ):
        """
```

Improved MOG2-based motion detector:

[IPCVC Syllabus Mapping - UNIT II: Image Segmentation and Alignment]

- Segmentation: Background subtraction via Mixture of Gaussians natively separates foreground (dynamic objects) from static background frames.
- Mode finding / Energy methods: Distinguishes shifting pixel energy values continuously to map valid contours.

- Gaussian blur to smooth noise before subtraction
- Dilation to fill small holes inside moving objects
- Erosion to separate nearby/touching blobs
- Area + aspect-ratio filtering to remove merged and spurious detections

```

"""
self.back_sub = cv2.createBackgroundSubtractorMOG2(
    history=history,
    varThreshold=var_threshold,
    detectShadows=detect_shadows,
)
self.min_area = min_area
self.max_area = max_area
self.min_ar = min_aspect_ratio
self.max_ar = max_aspect_ratio
self.learning_rate = learning_rate

# Morphology kernels kept as instance variables (avoid re-creating every frame)
self._open_kernel = cv2.getStructuringElement(cv2.MORPH_ELLIPSE, (3, 3))
self._close_kernel = cv2.getStructuringElement(cv2.MORPH_ELLIPSE, (9, 9))
self._dilate_kernel = cv2.getStructuringElement(cv2.MORPH_ELLIPSE, (5, 5))

```

```

# -----
# Internal helpers
# -----

```

```
def _preprocess(self, frame):
```

```

"""
    Apply Gaussian blur to the input frame before background subtraction.
    Reduces high-frequency sensor noise that causes spurious detections.
"""

```

```
    return cv2.GaussianBlur(frame, (5, 5), 0)
```

```
def _postprocess_mask(self, mask):
```

```

"""
    Clean the raw foreground mask with a pipeline of morphological operations:
    1. OPEN (erode → dilate) – removes small isolated noise pixels
    2. CLOSE (dilate → erode) – fills holes inside object silhouettes
    3. Additional dilation – gently expands blobs so thin limbs merge
                                   back into the body silhouette
"""

```

```

mask = cv2.morphologyEx(mask, cv2.MORPH_OPEN, self._open_kernel)
mask = cv2.morphologyEx(mask, cv2.MORPH_CLOSE, self._close_kernel)
mask = cv2.dilate(mask, self._dilate_kernel, iterations=1)

```

```

return mask

def _is_valid_detection(self, x, y, w, h):
    """
    Returns True when the bounding box passes all sanity checks:
    • Area within [min_area, max_area]
    • Aspect ratio (w/h) within [min_ar, max_ar]
    """
    area = w * h
    if area < self.min_area or area > self.max_area:
        return False
    ar = w / float(h) if h > 0 else 0
    if ar < self.min_ar or ar > self.max_ar:
        return False
    return True

# -----
# Public API
# -----

def apply(self, frame):
    """
    Full detection pipeline.

    Returns:
    detections : list of ([x, y, w, h], confidence, class_label)
                  ready for DeepSORT's update_tracks().
    fg_mask    : cleaned binary foreground mask (for debug visualisation).
    """
    # Step 1 – denoise the frame
    blurred = self._preprocess(frame)

    # Step 2 – background subtraction with controlled learning rate
    raw_mask = self.back_sub.apply(blurred, learningRate=self.learning_rate)

    # Step 3 – threshold: keep only definite foreground (255), discard shadows (127)
    _, fg_mask = cv2.threshold(raw_mask, 254, 255, cv2.THRESH_BINARY)

    # Step 4 – clean the mask
    fg_mask = self._postprocess_mask(fg_mask)

    # Step 5 – find external contours
    contours, _ = cv2.findContours(
        fg_mask, cv2.RETR_EXTERNAL, cv2.CHAIN_APPROX_SIMPLE
    )

    # Step 6 – build DeepSORT detection list with quality filtering
    detections = []
    for cnt in contours:
        x, y, w, h = cv2.boundingRect(cnt)

```

```

if not self._is_valid_detection(x, y, w, h):
    continue

# Confidence proportional to filled ratio (contour area / bbox area)
# Penalises very sparse / ghostly contours
filled_ratio = cv2.contourArea(cnt) / float(w * h)
confidence = float(np.clip(filled_ratio, 0.3, 1.0))

# DeepSORT expects: ([left, top, width, height], confidence, class)
detections.append([x, y, w, h], confidence, "object")

return detections, fg_mask

```

--- preprocessing.py ---

"""

Step 2: Image Preprocessing Pipeline

=====

Reads extracted frames from dataset_frames/, applies resize + denoise + CLAHE, saves to processed_dataset/, and generates before/after samples.

Usage:

```

python src/preprocessing.py
python src/preprocessing.py --size 224 --samples 10
"""

```

```

import os
import sys
import glob
import random
import argparse
import cv2
import numpy as np
from tqdm import tqdm

```

```

_SRC_DIR = os.path.dirname(os.path.abspath(__file__))
if _SRC_DIR not in sys.path:
    sys.path.insert(0, _SRC_DIR)

```

```

BASE_DIR = os.path.dirname(_SRC_DIR)

```

```

def preprocess_image(image, target_size=224):
    """

```

Apply the full preprocessing pipeline to a single image.

[IPCV Syllabus Mapping - UNIT I: Image Processing]

- Geometric Transformations: Resizing pixel matrices.
- Linear Filtering: Gaussian kernels to convolute matrices.
- Point Operators: CLAHE performing non-linear contrast adjustments.

Returns:

```

    processed : preprocessed BGR image
    """
    # 1. Geometric Transformation: Resize (Affine bounds)
    resized = cv2.resize(image, (target_size, target_size))

    # 2. Linear Filtering: Gaussian Denoise (5x5 spatial domain)
    # Suppresses high-frequency artifacts safely.
    denoised = cv2.GaussianBlur(resized, (5, 5), 0)

    # 3. Point Operators: CLAHE on luminance (L channel)
    lab = cv2.cvtColor(denoised, cv2.COLOR_BGR2LAB)
    l_chan, a_chan, b_chan = cv2.split(lab)
    clahe = cv2.createCLAHE(clipLimit=2.0, tileGridSize=(8, 8))
    l_enhanced = clahe.apply(l_chan)
    lab_enhanced = cv2.merge([l_enhanced, a_chan, b_chan])
    processed = cv2.cvtColor(lab_enhanced, cv2.COLOR_LAB2BGR)

    return processed

```

```

def make_before_after(before, after, target_h=300):
    """Create a side-by-side comparison image with labels."""
    # Resize both to same height for clean display
    h1, w1 = before.shape[:2]
    h2, w2 = after.shape[:2]
    before_disp = cv2.resize(before, (int(w1 * target_h / h1), target_h))
    after_disp = cv2.resize(after, (int(w2 * target_h / h2), target_h))

    # Add labels
    cv2.putText(before_disp, "BEFORE", (10, 30),
                cv2.FONT_HERSHEY_SIMPLEX, 0.8, (0, 0, 255), 2)
    cv2.putText(after_disp, "AFTER", (10, 30),
                cv2.FONT_HERSHEY_SIMPLEX, 0.8, (0, 255, 0), 2)

    # Stack horizontally
    combined = np.hstack([before_disp, after_disp])
    return combined

```

```

def run_preprocessing(
    input_dir,
    output_dir,
    sample_dir,
    target_size=224,
    num_samples=10,
):
    """Main preprocessing pipeline."""

    categories = ["person", "vehicle", "background"]
    all_pairs = [] # (before_path, after_path) for sampling

```

```

for cat in categories:
    src_folder = os.path.join(input_dir, cat)
    dst_folder = os.path.join(output_dir, cat)
    os.makedirs(dst_folder, exist_ok=True)

    images = sorted(glob.glob(os.path.join(src_folder, "*.jpg")))
    if not images:
        print(f" [{cat}] No images found, skipping.")
        continue

    print(f" [{cat}] Processing {len(images)} images...")

    for img_path in tqdm(images, desc=f" {cat}", unit="img", leave=False):
        img = cv2.imread(img_path)
        if img is None:
            continue

        processed = preprocess_image(img, target_size=target_size)

        fname = os.path.basename(img_path)
        out_path = os.path.join(dst_folder, fname)
        cv2.imwrite(out_path, processed, [cv2.IMWRITE_JPEG_QUALITY, 95])

        all_pairs.append((img_path, out_path))

# ——— Generate before/after samples ———
os.makedirs(sample_dir, exist_ok=True)
num_samples = min(num_samples, len(all_pairs))

if num_samples > 0:
    samples = random.sample(all_pairs, num_samples)
    print(f"\n Generating {num_samples} before/after samples...")

    for i, (before_path, after_path) in enumerate(samples):
        before_img = cv2.imread(before_path)
        after_img = cv2.imread(after_path)
        if before_img is None or after_img is None:
            continue

        comparison = make_before_after(before_img, after_img)
        save_path = os.path.join(sample_dir, f"sample_{i+1:02d}.jpg")
        cv2.imwrite(save_path, comparison, [cv2.IMWRITE_JPEG_QUALITY, 95])

total = len(all_pairs)
print(f"\n[Preprocessing] Done! {total} images processed.")
print(f" Output: {output_dir}")
print(f" Samples: {sample_dir}")
return total

```

```
def main():
    parser = argparse.ArgumentParser(description="Image Preprocessing Pipeline")
    parser.add_argument("--input", type=str,
                        default=os.path.join(BASE_DIR, "dataset_frames"),
                        help="Input folder (from data_preparation.py)")
    parser.add_argument("--output", type=str,
                        default=os.path.join(BASE_DIR, "processed_dataset"),
                        help="Output folder for preprocessed images")
    parser.add_argument("--samples-dir", type=str,
                        default=os.path.join(BASE_DIR, "sample_before_after"),
                        help="Output folder for before/after comparisons")
    parser.add_argument("--size", type=int, default=224,
                        help="Target image size (224 or 640)")
    parser.add_argument("--samples", type=int, default=10,
                        help="Number of before/after samples to generate")
    args = parser.parse_args()

    run_preprocessing(
        input_dir=args.input,
        output_dir=args.output,
        sample_dir=args.samples_dir,
        target_size=args.size,
        num_samples=args.samples,
    )

if __name__ == "__main__":
    main()

--- trajectory.py ---

import collections
import cv2

class TrajectoryTracker:
    def __init__(self, max_points=50):
        """
        Initializes the Trajectory Tracker.

        [IPCV Syllabus Mapping - UNIT V: Object Measurement, Pose Estimation and Understanding]
        - Object Measurement: Counting continuous positional locations to map speed vectors and
          trajectory lengths.
        - Image Understanding: Connecting spatial dots to visually output analytical path trails.

        max_points: Max number of history points to keep for a trajectory.
        """
        # Dictionary to store trajectories. Key is track_id, value is a deque of points.
        self.trajectories = collections.defaultdict(lambda: collections.deque(maxlen=max_points))
        self.max_points = max_points

    def update(self, track_id, center_point):
```

```

"""
Adds a new center point to the trajectory of track_id.
"""
self.trajectories[track_id].append(center_point)

def draw(self, frame, track_id, color=(0, 255, 255), thickness=2):
    """
    Draws the trajectory lines for a specific track_id on the given frame.
    """
    if track_id not in self.trajectories:
        return

    points = list(self.trajectories[track_id])
    for i in range(1, len(points)):
        if points[i - 1] is None or points[i] is None:
            continue

        # Draw line between previous point and current point
        cv2.line(frame, points[i - 1], points[i], color, thickness)

def cleanup(self, active_track_ids):
    """
    Removes trajectories for objects that have disappeared (not in active_track_ids).
    Helps keep memory usage efficient.
    """
    # Find track IDs to remove
    ids_to_remove = [tid for tid in self.trajectories.keys() if tid not in active_track_ids]
    for tid in ids_to_remove:
        self.trajectories.pop(tid, None)

```

APPENDIX B - SAMPLE OUTPUT SCREENSHOTS

